


VM Live Migration




- Migrate an entire VM from one physical host to another
 - All user processes and kernel state
 - Without having to shut down the machine
- Why migrate VMs?
 - Distribute VM load efficiently across servers in a cloud
 - System maintenance
- Easier than migrating processes
 - VM has a much narrower interface than a process
- Two main techniques: pre-copy and post-copy

“Live Migration of Virtual Machines”, Christopher Clark, Keir Fraser, Steven Hand, Jacob Gorm Hansen, Eric Jul, Christian Limpach, Ian Pratt, Andrew Warfield

“Post-Copy Based Live Virtual Machine Migration Using Adaptive Pre-Paging and Dynamic Self-Ballooning”, Michael R. Hines and Kartik Gopalan

What is migrated?



- CPU context of VM, contents of main memory
 - Narrow interface, easier than process migration
- Disk: assume NAS (network attached storage) that is accessible from both hosts, or local disk is mirrored
 - We do not consider migrating disk data
- Network: assume both hosts on same LAN
 - Migrate IP address, advertise new MAC address to IP mapping via ARP reply
 - Migrate MAC address, let switches learn new MAC location
 - Network packets redirected to new location (with transient losses)
- I/O devices are provisioned at target
 - Virtual I/O devices easier to migrate, direct device assignment of physical devices to VMs (device passthrough) makes migration harder

Steps to migrate a VM

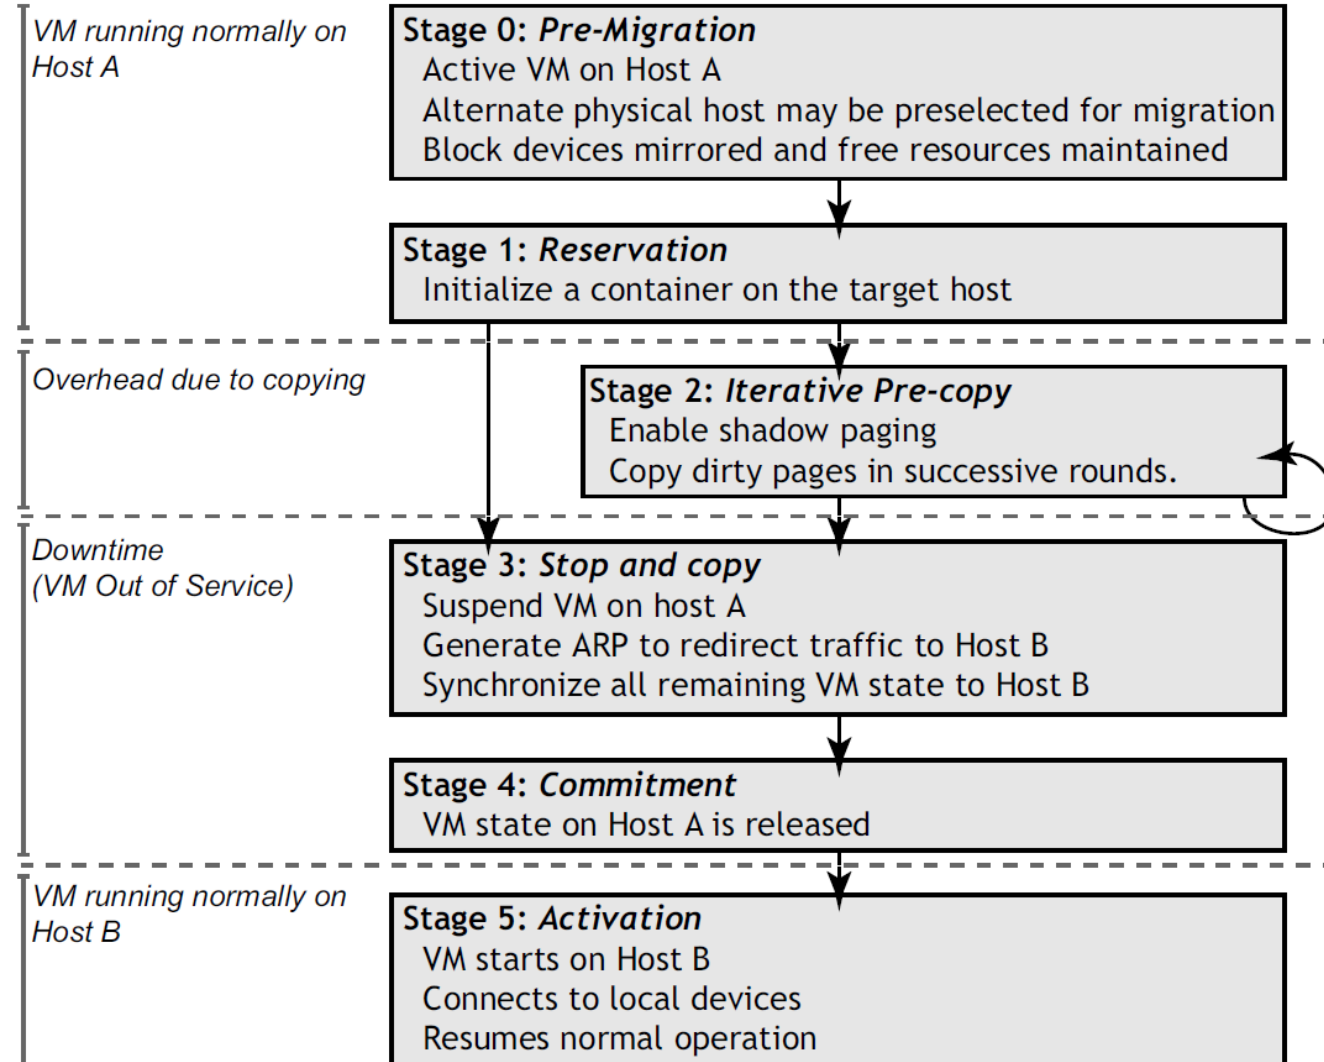
- Broad steps in any migration technique: Suppose we are migrating a VM from host A to host B
 1. Setup target host B, reserve resources for the VM
 2. **Push phase**: push some memory of VM from A to B
 3. **Stop-and-copy**: stop the VM at A, copy CPU context, and some memory
 4. **Pull phase**: Start VM at host B, pull any further memory required from A
 5. Clean up state from host A, migration complete
- **Total migration time**: time for steps 2,3,4
- **Service downtime**: time for step 3
- Other metrics: impact on application performance, network bandwidth consumed, total pages transferred

Flavors of migration techniques


- **Pure stop-and-copy**: VM stopped, all state transferred to target, VM restarts
 - Too much downtime to be classified as “live” migration
- **Pre-copy**: most state is transferred in the push phase, followed by a brief stop-and-copy phase
- **Post-copy**: VM stopped, bare minimum state required to run the VM is transferred to the target host. Remaining state is pulled on demand while the VM is running at the new location.
- **Hybrid**: a mix of pre-copy and post-copy. Some pushing of state followed by stop-and-copy, followed by pulling of state on demand.

Pre-copy based live migration

- Iterative pre-copy + stop-and-copy for remaining memory
- First push round copies all pages
- Every round copies pages dirtied in previous round
 - A page maybe copied multiple times
- Writable Working Set (WWS): pages commonly written to
 - WWS will be copied multiple times
 - Finally transferred in stop-and-copy
- How many rounds? Stop when rate of dirtying > rate of transfer
 - Diminishing returns with more than few rounds




Impact of iterative pre-copy




- If stop-and-copy, 512MB VM, 128 Mbps network, downtime = 32 sec
- With one pre-copy round, downtime goes to 2-3 sec
 - ~1 second for 2 or more rounds

Tracking dirty pages



- Xen-based implementation
 - Page tables in Xen maintained by guest
 - Move to **shadow page tables** for migration
 - Migration managed by control software in domain0
- Shadow page table constructed on demand for every round
 - **Dirty bitmap** maintained for every round
 - Any page access by guest → page fault to Xen, shadow page table updated
 - PTE marked as read-only by default in shadow
 - If valid write access, shadow PTE marked writeable, page marked dirty in bitmap
 - At end of round, dirty pages are marked for transfer in control software
 - Shadow page table and dirty bitmap reinitialized after every round
 - Last set of dirty pages copied in stop-and-copy
- Guest page table in target host changed based on new physical addresses

Shadow page table in Xen




Some optimizations

- Avoid transferring page multiple times
 - Before transmitting page, peek into the current round's dirty bitmap
 - Skip transmission if page is already dirtied in ongoing round
- Move non-interactive processes generating dirty pages to wait queue
 - Execution paused until migration completes
- Free up page cache and other unnecessary pages
 - Reduce memory footprint
 - Much like ballooning


Pre-copy performance

- Downtime: ~100 millisecc, total migration time of few tens of seconds
- Worse for memory-intensive applications, better for interactive apps

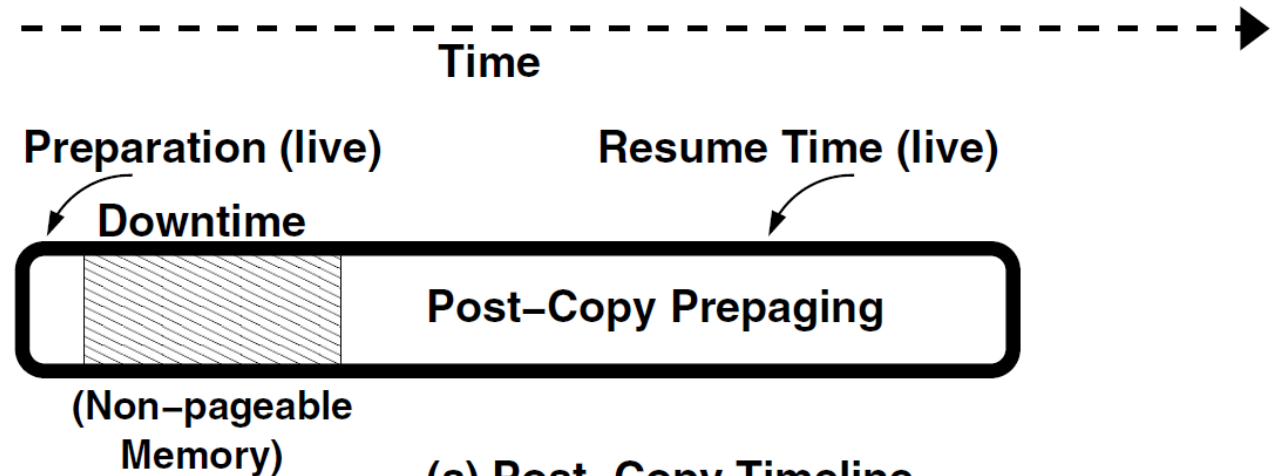


Post-copy based live migration

- Avoid multiple transfers of same page as happens in pre-copy
- Prepare target, stop VM, copy CPU context and minimum memory to target
- Start VM at target, pull memory from source via demand paging
 - Memory access at target causes page fault, page fetched from source



(a) Pre-Copy Timeline

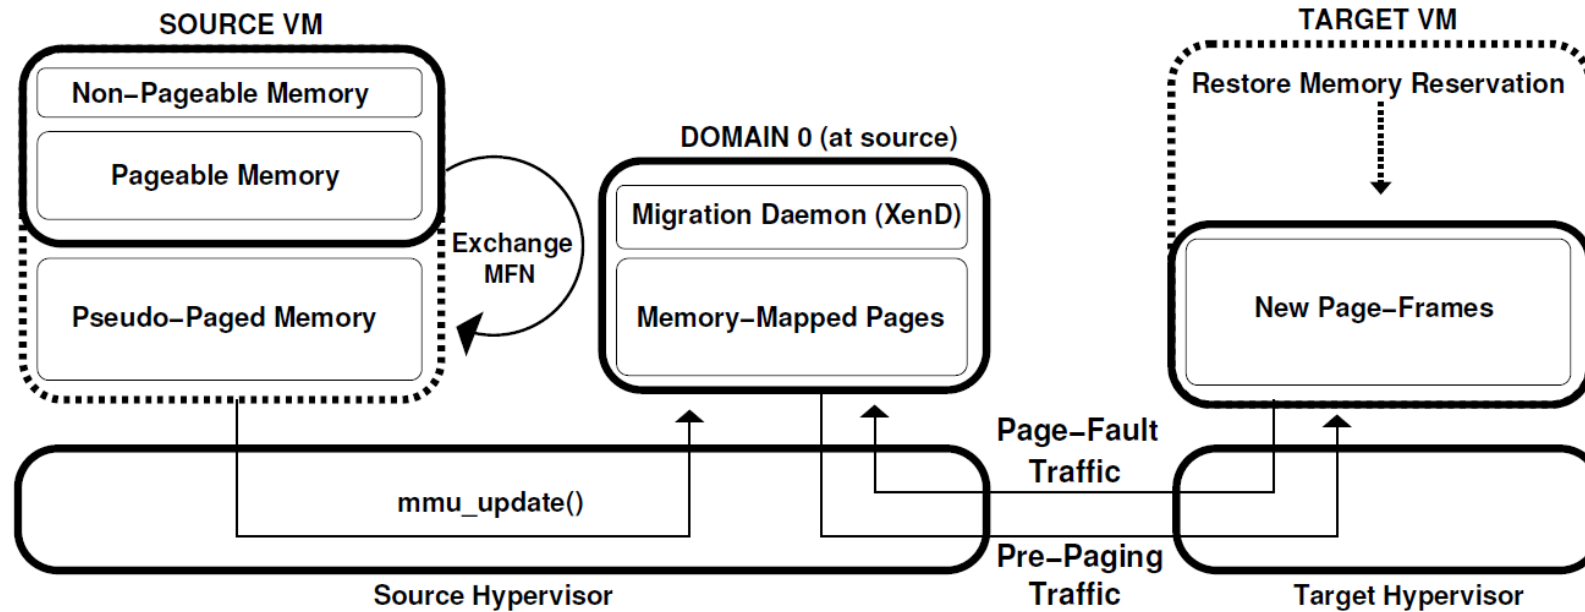


(a) Post-Copy Timeline

Optimizations

- **Active pushing**: source proactively pushes important pages, in addition to pulling pages via page faults
- **Pre-paging**: a “bubble” of pages around faulted page and proactively pushed, in anticipation of future accesses
- **Dynamic self-ballooning**: VM periodically frees up unnecessary memory and gives it back to hypervisor
 - Reduces memory footprint, speeds up page transfer
 - Performed carefully without hurting application performance
 - Can be used to optimize pre-copy migration as well
- **Hybrid**: one pre-copy round, followed by post copy

Implementation details (Xen)



- How are pages pulled at target? “Pseudo-paging”
 - Page to a pseudo, in-memory, swap device (part of domain0). No memory copy, just transfer pages across domains. Guest page table updated suitably.
 - Only non-pageable memory transferred during stop-and-copy
 - When guest resumes at target, fetch memory from pseudo-paging device via page fault mechanism
 - Special swap device driver fetches from source over the network
- Alternative: use shadow page tables
 - If page fault to non-existent page at target, trap to hypervisor, fetch from source and update

What about failures?

- What if target machine fails during migration?
- Pre-copy can simply abort the migration, restart with another target
 - With pre-copy, latest state is on source only, so can recover
- With post copy, source has stale memory, target has updated memory
 - If target crashes during post copy, cannot recover application data (unless some replication is performed)

Post copy performance

- Longer downtime as compared to pre-copy, but lower total migration time, fewer page transfers, lesser disruption to application




Figure 9. Impact of post-copy on NetPerf bandwidth.


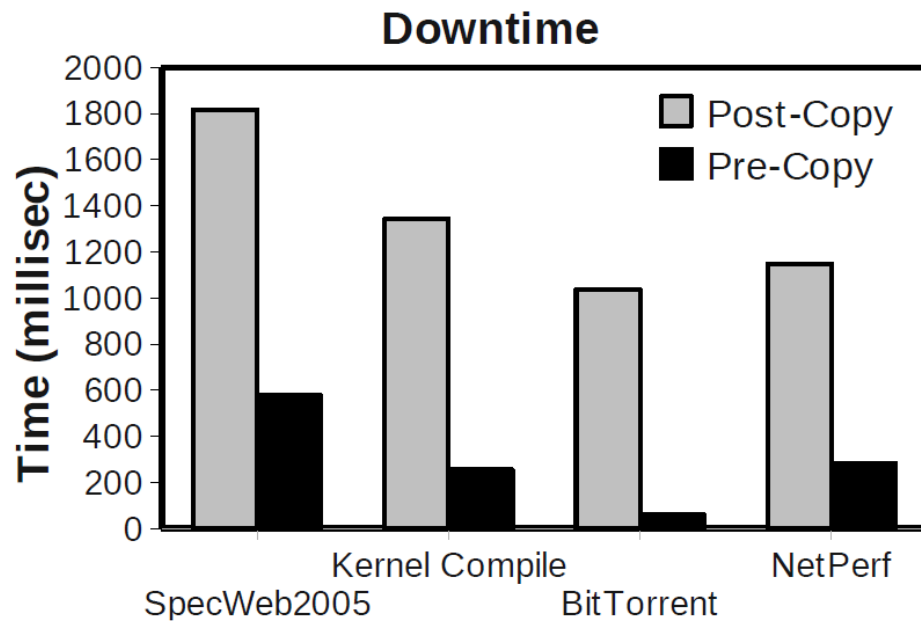
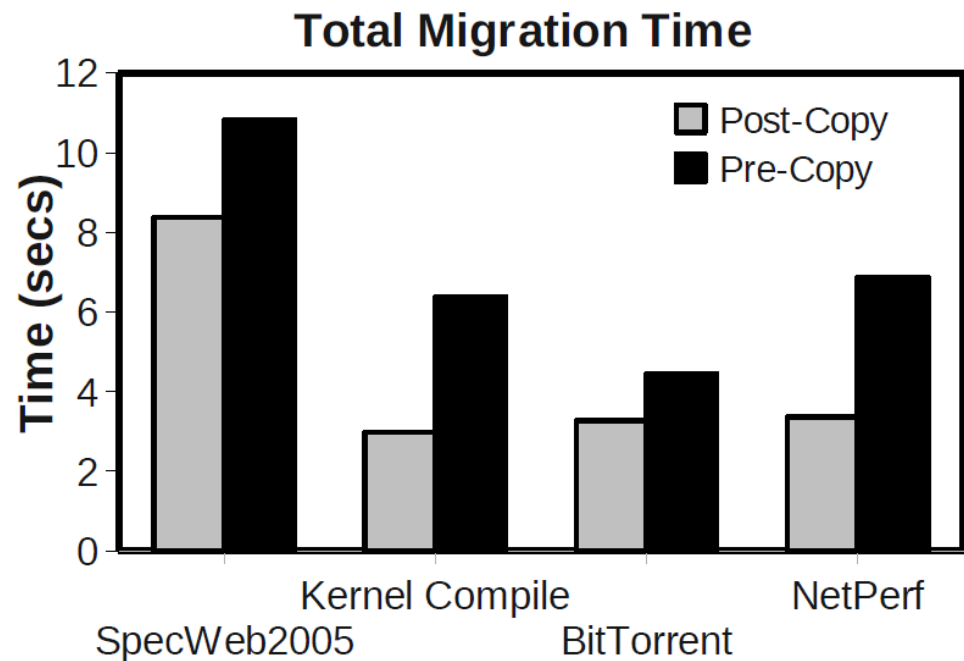
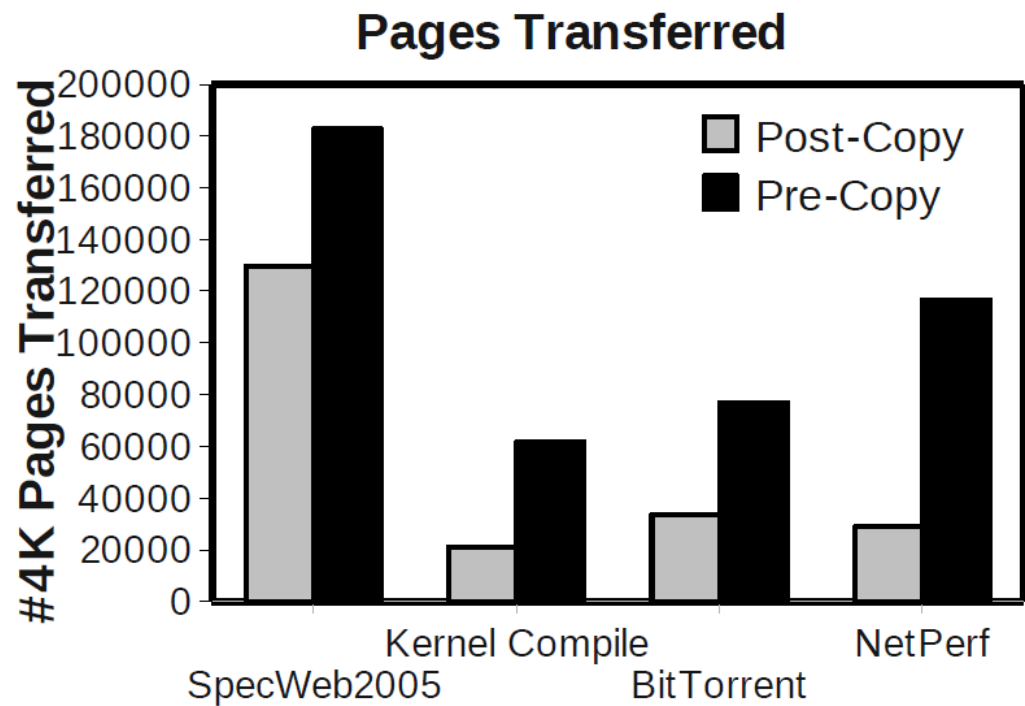


Figure 10. Impact of pre-copy on NetPerf bandwidth.



Summary

- VM live migration techniques
 - Iterative pre-copy vs post-copy via demand paging
 - Implementation details on Xen
 - Performance comparison
- Which is better?
 - Pre-copy suited for interactive application
 - Post copy is better for memory-intensive applications with large WWS
 - Hybrid techniques are also used